

Marcus Vance

marcus.vance@example.com | +1 (555) 876-5432 | Austin, TX | github.com/marcusvance | linkedin.com/in/marcusvance

PROFESSIONAL SUMMARY

Backend & Machine Learning Engineer with 4+ years specializing in distributed systems, high-throughput microservices, and Generative AI integrations. Skilled in Python, FastAPI, PostgreSQL, Docker, Kubernetes, and LLM orchestration with LangChain.

TECHNICAL SKILLS

Core Competencies: Python, FastAPI, Django, Node.js, PostgreSQL, Redis, MongoDB, Docker, Kubernetes, AWS (EC2, S3, RDS, Lambda), LangChain, OpenAI API, Gemini API, PyTorch, gRPC, Kafka, Linux, Terraform, PyTest

PROFESSIONAL EXPERIENCE

Staff Backend & AI Engineer — *Cognitive Cloud Solutions* | 2022 - Present

- Developed scalable FastAPI microservices integrating Google Gemini and OpenAI models with semantic retrieval (RAG).
- Engineered PostgreSQL database schemas and optimized indexing, reducing P99 query latency by 55%.
- Deployed containerized Kubernetes clusters on AWS EKS with auto-scaling to handle 15,000 requests per minute.
- Implemented Redis caching layer reducing external LLM token expenditure by 30% month-over-month.

Software Engineer (Backend) — *DataMatrix Systems* | 2020 - 2022

- Designed REST APIs in Python/Django for enterprise workflow analytics platform.
- Created asynchronous background job queues using Celery and Redis to process 500K daily batch records.
- Authored comprehensive unit and integration test suites using PyTest maintaining 92% coverage.

EDUCATION & CERTIFICATIONS

B.S. in Software Engineering — University of Texas at Austin (2020)

AWS Certified Solutions Architect – Associate — Amazon Web Services (2023)